

# Deep Learning for Computer Vision HW#3

B05901182 電機四 潘彥銘

## Problem 1: (20%)

Collaboration: None

Reference:

[https://pytorch.org/tutorials/beginner/dcgan\\_faces\\_tutorial.html](https://pytorch.org/tutorials/beginner/dcgan_faces_tutorial.html)

### 1. (5%)

The structure of the generator:

```
G: Generator(
  (layer1): Sequential(
    (0): ConvTranspose2d(100, 512, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace)
  )
  (layer2): Sequential(
    (0): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace)
  )
  (layer3): Sequential(
    (0): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace)
  )
  (layer4): Sequential(
    (0): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace)
  )
  (layer5): Sequential(
    (0): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): Tanh()
  )
)
```

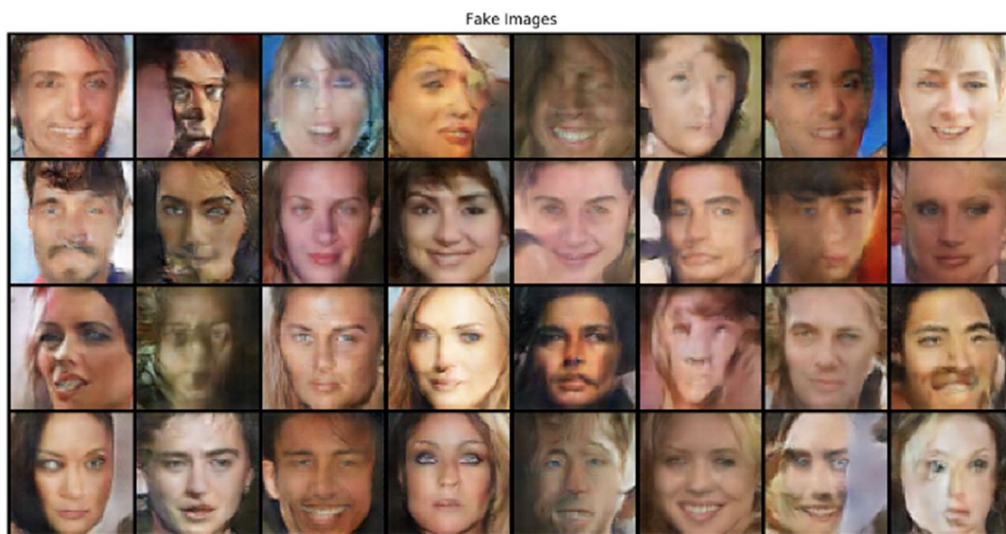
The structure of the discriminator:

```
D: Discriminator(
  (layer1): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace)
  )
  (layer2): Sequential(
    (0): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): LeakyReLU(negative_slope=0.2, inplace)
  )
  (layer3): Sequential(
    (0): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): LeakyReLU(negative_slope=0.2, inplace)
  )
  (layer4): Sequential(
    (0): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): LeakyReLU(negative_slope=0.2, inplace)
  )
  (layer5): Sequential(
    (0): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): Sigmoid()
  )
)
```

The other hyper-parameters of the my model:

```
n_workers = 4
batch_size = 128
image_size = 64
num_channels = 3
z_size = 100 # Size of z latent vector (i.e. size of generator input)
gf_size = 64 # Size of feature maps in generator
df_size = 64 # Size of feature maps in discriminator
num_epochs = 50
lr = 0.0002
beta1 = 0.5
beta2 = 0.999 # (beta1, beta2) are the hyper-parameters of Adam optim
num_gpus = 1
```

2. (10%)



3. (5%)

While I implemented my GAN model, I was confused that why I should add `.detach()` while training discriminator (like the figure below).

```
output = netD(fake.detach()).view(-1)
```

After looking up it on google, I found that this is really important. While training discriminator, we don't want to update generator. So the `.detach()` is necessary. If we do not detach it, then the backward process will clear all the variable on the graph, including fake. Thus generator won't be update next time when we call fake.

## Problem 2: (20%)

Collaboration: B05901027 詹書愷

Reference:

1. <https://arxiv.org/abs/1610.09585>
2. <https://github.com/clvr/ai/ACGAN-PyTorch>

### 1. (5%)

The structure of the generator:

```
G: Generator(
  (layer0): Sequential(
    (0): Linear(in_features=101, out_features=384, bias=True)
  )
  (layer1): Sequential(
    (0): ConvTranspose2d(384, 192, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace)
  )
  (layer2): Sequential(
    (0): ConvTranspose2d(192, 96, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace)
  )
  (layer3): Sequential(
    (0): ConvTranspose2d(96, 48, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace)
  )
  (layer4): Sequential(
    (0): ConvTranspose2d(48, 24, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace)
  )
  (layer5): Sequential(
    (0): ConvTranspose2d(24, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): Tanh()
  )
)
```

The structure of the discriminator:

```
D: Discriminator(
  (layer1): Sequential(
    (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace)
    (2): Dropout(p=0.5)
  )
  (layer2): Sequential(
    (0): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), bias=False)
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): LeakyReLU(negative_slope=0.2, inplace)
    (3): Dropout(p=0.5)
  )
  (layer3): Sequential(
    (0): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): LeakyReLU(negative_slope=0.2, inplace)
    (3): Dropout(p=0.5)
  )
  (layer4): Sequential(
    (0): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), bias=False)
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): LeakyReLU(negative_slope=0.2, inplace)
    (3): Dropout(p=0.5)
  )
  (layer5): Sequential(
    (0): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): LeakyReLU(negative_slope=0.2, inplace)
    (3): Dropout(p=0.5)
  )
  (layer6): Sequential(
    (0): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): LeakyReLU(negative_slope=0.2, inplace)
    (3): Dropout(p=0.5)
  )
  (dis_fc): Linear(in_features=12800, out_features=1, bias=True)
  (ac_fc): Linear(in_features=12800, out_features=2, bias=True)
  (sigmoid): Sigmoid()
)
```

The other hyper-parameters of the my model:

```
n_workers = 4
batch_size = 100
image_size = 64
num_channels = 3
num_classes = 2
z_size = 100 # Size of z latent vector (i.e. size of generator input)
num_epochs = 100
lr = 0.0002
beta1 = 0.5
beta2 = 0.999 # (beta1, beta2) are the hyper-parameters of Adam optim
num_gpus = 1
```

2. (10%)



3. (5%)

At first, my generator's output image doesn't smile when I feed a noise vector concatenated with a smile label. After I checked my code several times, I found that I made a mistake.

When I'm training the auxiliary classifier along with generator, I feed the true label to classifier, which is not right. I should feed the fake label, which is randomly generated, to the auxiliary classifier. Because now the ac\_output (like the figure below) is generated by discriminator, which is feeded by the fake image. So the auxiliary classifier should be feeded the fake label to learn correctly.

```
dis_output, ac_output = netD(fake) |
Gloss = dis_criterion(dis_output, dis_label) + ac_criterion(ac_output, fake_smile_label)
```

### Problem 3: (35%)

Collaboration:

B05901027 詹書愷、B05602042 林奕廷、B05901074 陳泓均

Reference:

1. <http://sites.skoltech.ru/compvision/projects/grl/files/paper.pdf>
2. [https://github.com/fungtion/DANN\\_py3](https://github.com/fungtion/DANN_py3)


1.(3%) 、 2.(10%) 、 3. (3%)

	SVHN -> MNIST-M	MNIST-M -> SVHN
Trained on source	0.4523(45.23%)	0.3891(38.91%)
Adaptation(DANN)	0.4538(45.38%)	0.4484(44.84%)
Trained on target	0.9632(96.32%)	0.8842(88.42%)




4. (6%)


MNIST-M -> SVHN (a) different digit classes 0-9 (I use all test data)




MNIST-M -> SVHN (b) different domains (I use all test data)



SVHN -> MNIST-M (a) different digit classes 0-9 (I use all test data)



SVHN -> MNIST-M (b) different domains (I use part of the data)



## 5. (6%)

The architecture of the two domain adaptation DANN model:

```
DANN(  
  (feature_extractor): Sequential(  
    (0): Conv2d(3, 32, kernel_size=(5, 5), stride=(1, 1))  
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU(inplace)  
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (4): Conv2d(32, 48, kernel_size=(5, 5), stride=(1, 1))  
    (5): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (6): ReLU(inplace)  
    (7): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (class_clf): Sequential(  
    (0): Linear(in_features=768, out_features=100, bias=True)  
    (1): ReLU(inplace)  
    (2): Dropout(p=0.5)  
    (3): Linear(in_features=100, out_features=100, bias=True)  
    (4): BatchNorm1d(100, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ReLU(inplace)  
    (6): Linear(in_features=100, out_features=10, bias=True)  
  )  
  (domain_clf): Sequential(  
    (0): Linear(in_features=768, out_features=100, bias=True)  
    (1): ReLU(inplace)  
    (2): Linear(in_features=100, out_features=2, bias=True)  
  )  
)
```

The other hyper-parameters of the two domain adaptation DANN model:

```
n_workers = 4  
batch_size = 128  
num_channels = 3  
num_classes = 10  
num_epochs = 100  
lr = 0.0002  
num_gpus = 1
```

## 6. (7%)

At first, my domain adaptation DANN model's (SVHN -> MNIST-M) performance on accuracy is worse than the model (lower bound) which trained on source and tested on the target(43.06%<45.23%), which is really weird. Thus I checked my model again, and found that I added a dropout layer in the feature extractor part. My mind came up with an idea: The feature extractor is the dominant part of this model. If I add the dropout layer in this part, which may cause the model to learn incompletely. So I removed the dropout layer, and the accuracy result became 45.38%, which is slightly better than the lower bound.

I also found that the model which trained and tested on different domain



data would increase its accuracy while high epoches(~90). But the domain adaptation model reaches its accuracy peak in the early epoches(~30).

**Problem 4: (35%)**

Collaboration:

B05901027 詹書愷、B05602042 林奕廷、B05901074 陳泓均

Reference:


1. <https://arxiv.org/pdf/1702.05464.pdf>
2. <https://github.com/corenel/pytorch-adda>

1. (6+10%)


	SVHN -> MNIST-M	MNIST-M -> SVHN
Adaptation(ADDA)	0.5955(59.55%)	0.4581(45.81%)

2. (6%)


SVHN -> MNIST-M (a) different digits classes 0-9 (I use part of the data)




SVHN -> MNIST-M (b) different domains (I use part of the data)



MNIST-M -> SVHN (a) different digits classes 0-9 (I use part of the data)



MNIST-M -> SVHN (b) different domains (I use part of the data)



3. (6%)

My model has three types of sub-model: encoder, classifier and discriminator:

Encoder:

```
class Encoder(nn.Module):    # modified LeNet
    def __init__(self):
        super(Encoder, self).__init__()
        self.feature_extractor = nn.Sequential(
            nn.Conv2d(1, 20, kernel_size=5, stride=1),
            nn.MaxPool2d(kernel_size=2, stride=2, dilation=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(20, 50, kernel_size=5, stride=1),
            nn.Dropout2d(p=0.5),
            nn.MaxPool2d(kernel_size=2, stride=2, dilation=1),
            nn.ReLU(inplace=True)
        )
        self.fc1 = nn.Linear(800, 500)

    def forward(self, img):
        ft = self.feature_extractor(img)
        ft = ft.view(-1, 800)
        ft = self.fc1(ft)
        return ft
```

Classifier:

```
class Classifier(nn.Module):
    def __init__(self):
        super(Classifier, self).__init__()
        self.clf = nn.Sequential(
            nn.ReLU(),
            nn.Dropout(p=0.5),
            nn.Linear(500, 10)
        )

    def forward(self, ft):
        prob = self.clf(ft)
        return prob
```



Discriminator:

```
class Discriminator(nn.Module):
    def __init__(self):
        super(Discriminator, self).__init__()
        self.dis = nn.Sequential([
            nn.Linear(500, 500),
            nn.ReLU(inplace=True),
            nn.Linear(500, 2),
        ])

    def forward(self, img):
        prob = self.dis(img)

        return prob
```

The other hyper-parameters of the my model:

```
n_workers = 4
batch_size = 128
num_classes = 10
num_epochs = 100
lr = 0.0002
num_gpus = 1
beta1 = 0.5
```

I trained my source encoder and source classifier for 25 epochs, and initialize the weight of target encoder with the weight of target encoder. Then I trained the target encoder and discriminator for 100 epochs.

4. (7%)

Discuss what you've observed and learned from implementing your improved UDA model.

At first, my model's performance is really bad, the accuracy is between 9~20%, which is not better than the model of problem 3.

By observing the loss between classifier&encoder 、 discriminator, I found that is discriminator is too strong. Thus, I modified my discriminator model to a weaker version(delete one fc layer). And, after training my source encoder, I use the weight of source encoder to initialize my target encoder instead of doing nothing.

By doing so, the performance better better. The SVHN -> MNIST-M model's accuracy is 59%, which is really out of my expectation.